

Intro. To Machine Learning & Computer Vision




TABLE OF CONTENTS

01

Intro to Machine Learning

02

Introduction to Computer
Vision

03

Using NLP!





What is BigThink?

BigThink AI's goal is to bring awareness towards the growing field of artificial intelligence. We intend to advance our members' knowledge and experience with fields within and related to AI through meetings, discussions, seminars, and projects.

What our Club Does?

Pyoneers

Weekly “lectures” where
our leads do a brief
lecture followed by an
example

Edge

A group of students with
machine learning
research experience
work together on a
semester long research
project.





01

Intro to Machine Learning





What is Machine Learning?

- Machine learning is a subset field of artificial intelligence
- Utilizes mathematics and statistics to learn from data
- Before ML, we write programs with algorithms that have deterministic behavior
 - Ideal for scenarios with complex rule sets that are not deterministic
- Aims to generalize the patterns in the data
- Two Main Problems
 - Classification and Regression
- We start with training data that can be labeled or unlabeled.
- Then we feed into a machine learning model, evaluate, and then deploy

Ex (Spam detection): It is difficult to write all the rules to detect spam in emails. Machine learning can adapt to changes over time.





Supervised Learning

- Supervised learning is when the data is labeled, so we have the input and the expected output
- Algorithm predicts classes on new data

Examples

- k-Nearest Neighbors
- Linear Regression
- Decision Trees and Random Forests
- Neural Networks




Unsupervised Learning

- The data is not labeled
- Algorithm find patterns or groups on its own without knowing how they are grouped

Examples

- Clustering
- Anomaly Detection




Reinforcement Learning

- Between supervised and unsupervised
- Does not rely on labeled data, but performs “desired” actions based on rewards
- Key Terminology:
 - Agent - What moves
 - State - Where the agent moves
 - Action - A move that can be taken
- The system gives a reward based on its actions, and updates what action it should take to maximize reward

Examples

- Games!
- Self driving cars!




What are the state, action, and agent in a game of chess?



02

Intro to Computer Vision





What is Computer Vision (CV)?

- Computer vision detects patterns and objects in images by identify each pixel to learn patterns
- New developments in Neural Networks paved way for computer vision
- Common applications are image segmentation, classification, object tracking, and detection



How does it work?

- We are given photos. The model will analyze the photos by each pixel
- Each pixel has 3 channels (for RGB) that ranges from 0-255 in red, green, and blue channels
- We then flatten the image to a single vector to feed into our model
- The model then learns to group or categorize the images based on similarities it finds in the patterns of these pixels, without being explicitly told what to look for




Image Array

$$\begin{bmatrix} 010 \\ 010 \\ 010 \end{bmatrix}$$


With flatten operation

$$[010\ 010\ 010]$$



Modern Challenges in CV



- Weak performance to train our models
- Training on large images is computationally expensive
- Computer vision (CV) models are sensitive to noise and cannot generalize patterns
- Poor data quality
- Not enough data





Solutions (Weak performance to train our models)


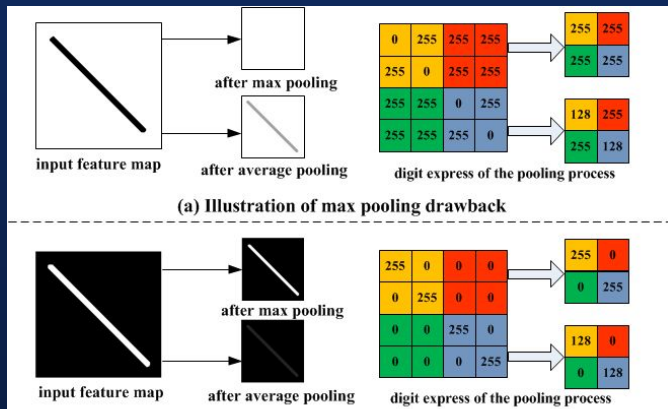
- Originally CV models ran on CPUs which was not enough to train our complex and large models
- Advancement to GPUs lead to improved calculations for CV models!
- Availability increased and costs decreased





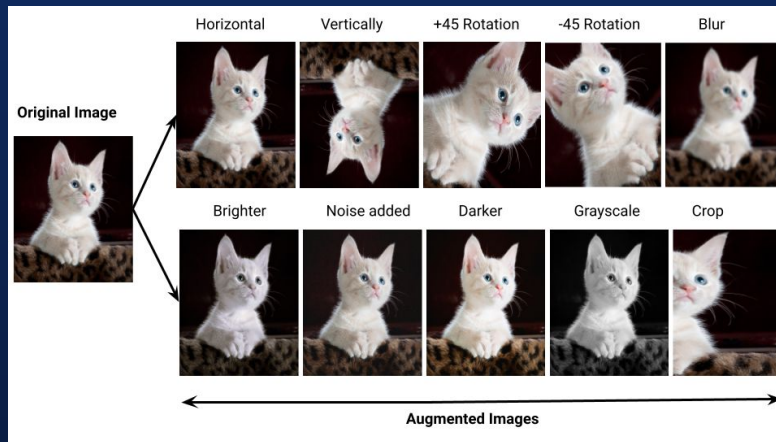
Solutions (Sensitive to Noise/Computation)

- Our models frequently overgeneralize (overfit) and do not learn features
- The images are large → increased computation
 - Solution: Pooling!
 - Each colored box below is a window. We summarize the features in that window by an operation
 - Max pooling (Take the max RGB value in a group of $n \times n$ pixels)



Solutions (Data Collection)

- A lack of images, results in the model failing to recognize patterns and poor performance
 - Solution: Data Augmentation!
 - We take the original data and modify the image to create new data to feed into the model
 - Methods: rotate, grayscale, color jitter, blur, invert, flip the image (There are more methods!)





03


RL Example





We will be doing, an Reinforcement learning example!

- **Packages**
 - random
 - gym
 - tensorflow
 - rl
- Use `pip install _` where blank is one of the above package names




Requirements

- **Problem:** A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The pendulum is placed upright on the cart and the goal is to balance the pole by applying forces in the left and right direction on the cart.

State: The space with information on the cart pole

Actions: Push Cart Left, or Push Cart Right

Agent: The decision-maker that tries to solve the task of balancing the pole





Time to Code!



THANK YOU!

